

ASSINMENT 11.1

1) STEP 1: Create database db2.

```
mysql> create database db2;
Query OK, 1 row affected (0.05 sec)

mysql> show databases;
+-----+
| Database |
+-----+
| information_schema |
| db |
| db2 |
| metastore |
| mysql |
+-----+
5 rows in set (0.11 sec)
```

2) Use database

```
mysql> use db2;
Database changed
mysql> █
```

3) Create table company.

```
mysql> create table company
-> (
-> id int,
-> name varchar(20),
-> location varchar(20)
-> );
Query OK, 0 rows affected (0.11 sec)

mysql> █
```

4) insert data into table

```
mysql> select * from company;
Empty set (0.00 sec)

mysql> insert into company values(1, 'acadgild','Bangalore');
Query OK, 1 row affected (0.02 sec)

mysql>
mysql> insert into company values(2, 'eduvista','Walnut');
Query OK, 1 row affected (0.00 sec)

mysql>
mysql> insert into company values(3, 'eduvista','Bangalore');
Query OK, 1 row affected (0.00 sec)

mysql> select * from company;
+-----+-----+-----+
| id    | name    | location |
+-----+-----+-----+
| 1     | acadgild | Bangalore |
| 2     | eduvista | Walnut    |
| 3     | eduvista | Bangalore |
+-----+-----+-----+
3 rows in set (0.00 sec)

mysql> █
```

4) Import data from mysql to hdfs :

```
acacgild@localhost ~$ sqoop import --connect jdbc:mysql://localhost/db2 \
--username 'root' -P --table 'company' --target-dir '/sqoopout' \
-m 1;
warning: /home/acacgild/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
warning: /home/acacgild/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
warning: /home/acacgild/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./zookeeper does not exist! Accumulo imports will fail.
Please set $ZOOKEEPER_HOME to the root of your Zookeeper installation.
2017-12-03 14:08:05,642 INFO [main] sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
2017-12-03 14:08:14,269 INFO [main] manager.MySQLManager: Preparing to use a MySQL streaming resultset.
2017-12-03 14:08:14,269 INFO [main] tool.CodeGenTool: Beginning code generation
2017-12-03 14:08:15,674 INFO [main] manager.SqlManager: Executing SQL statement: SELECT t.* FROM `company` AS t LIMIT 1
2017-12-03 14:08:15,799 INFO [main] manager.SqlManager: Executing SQL statement: SELECT t.* FROM `company` AS t LIMIT 1
2017-12-03 14:08:15,888 INFO [main] orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acacgild/hadoop-2.7.2
note: /tmp/sqoop-acacgild/compile/390c456869fc7e274f48bede4b8c3544/Company.java uses or overrides a deprecated API.
note: Recompile with -Xlint:deprecation for details.
2017-12-03 14:08:35,710 INFO [main] orm.CompilationManager: Writing jar file: /tmp/sqoop-acacgild/compile/390c456869fc7e274f48bede4b8c3544/company.jar
2017-12-03 14:08:36,061 WARN [main] manager.MySQLManager: It looks like you are importing from mysql.
2017-12-03 14:08:36,061 WARN [main] manager.MySQLManager: This transfer can be faster! Use the --direct
2017-12-03 14:08:36,061 WARN [main] manager.MySQLManager: option to exercise a MySQL-specific fast path.
2017-12-03 14:08:36,061 INFO [main] manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)
2017-12-03 14:08:36,071 INFO [main] mapreduce.ImportJobBase: Beginning import of company
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acacgild/hbase-1.0.3/lib/slf4j-log4j12-1.7.7.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acacgild/hadoop-2.7.2/share/hadoop/common/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
Java HotSpot(TM) Client VM warning: You have loaded library /home/acacgild/hadoop-2.7.2/lib/native/libhadoop.so.1.0.0 which might have disabled stack guard. The VM will try to fix the stack guard now.
It's highly recommended that you fix the library with 'execstack -c <libfile>', or link it with '-z noexecstack'.
2017-12-03 14:08:36,975 WARN [main] util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java class
es where applicable
```

```

2017-12-03 14:09:51,413 INFO [main] mapreduce.Job: map 0% reduce 0%
2017-12-03 14:10:08,283 INFO [main] mapreduce.Job: map 100% reduce 0%
2017-12-03 14:10:09,312 INFO [main] mapreduce.Job: Job job_1512280424524_0001 completed successfully
2017-12-03 14:10:09,565 INFO [main] mapreduce.Job: Counters: 30
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=136483
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=87
    HDFS: Number of bytes written=60
    HDFS: Number of read operations=4
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Other local map tasks=1
    Total time spent by all maps in occupied slots (ms)=14805
    Total time spent by all reduces in occupied slots (ms)=0
    Total time spent by all map tasks (ms)=14805
    Total vcore-seconds taken by all map tasks=14805
    Total megabyte-seconds taken by all map tasks=15160320
  Map-Reduce Framework
    Map input records=3
    Map output records=3
    Input split bytes=87
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=577
    CPU time spent (ms)=1620
    Physical memory (bytes) snapshot=68366336
    Virtual memory (bytes) snapshot=323010560
    Total committed heap usage (bytes)=16318464
  File Input Format Counters
    Bytes Read=0
  File Output Format Counters

```

```

cable
Found 2 items
-rw-r--r--  1 acadgild supergroup          0 2017-12-03 14:10 /sqoopout/ SUCCESS
-rw-r--r--  1 acadgild supergroup        60 2017-12-03 14:10 /sqoopout/part-m-00000

```

Export data from hdfs to Mysql

```

mysql> use db2;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

```

```

Database changed
mysql>

```

```

Database changed
mysql> delete from company;
Query OK, 3 rows affected (0.02 sec)

mysql>

```

```

[acadgild@localhost ~]$ sqoop export --connect jdbc:mysql://localhost/db2 --username 'root' -P --table 'company' --export-dir '/sqoopout' --in
put-fields-terminated-by ',' -m 1 --columns id,name,location

```

```

warning: /home/acadgild/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
warning: /home/acadgild/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
warning: /home/acadgild/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./zookeeper does not exist! Accumulo imports will fail.
Please set $ZOOKEEPER_HOME to the root of your Zookeeper installation.
2017-12-04 11:41:03,352 INFO [main] sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
2017-12-04 11:41:08,022 INFO [main] manager.MySQLManager: Preparing to use a MySQL streaming resultset.
2017-12-04 11:41:08,023 INFO [main] tool.CodeGenTool: Beginning code generation
2017-12-04 11:41:08,681 INFO [main] manager.SqlManager: Executing SQL statement: SELECT t.* FROM `company` AS t LIMIT 1
2017-12-04 11:41:08,736 INFO [main] manager.SqlManager: Executing SQL statement: SELECT t.* FROM `company` AS t LIMIT 1
2017-12-04 11:41:08,759 INFO [main] orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/hadoop-2.7.2
Note: /tmp/sqoop-acadgild/compile/dd01f8638f6beb10c1a97cd7b40a9219/company.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
2017-12-04 11:41:12,214 INFO [main] orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/dd01f8638f6beb10c1a97cd7b40a9219/company.jar
2017-12-04 11:41:12,422 INFO [main] mapreduce.ExportJobBase: Beginning export of company
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/hbase-1.0.3/lib/slf4j-log4j12-1.7.7.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/hadoop-2.7.2/share/hadoop/common/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
Java HotSpot(TM) Client VM warning: You have loaded library /home/acadgild/hadoop-2.7.2/lib/native/libhadoop.so.1.0.0 which might have disabled stack guard. The VM will try to fix the stack guard now.
It's highly recommended that you fix the library with 'execstack -c <libfile>', or link it with '-z noexecstack'.
2017-12-04 11:41:12,914 WARN [main] util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
2017-12-04 11:41:12,938 INFO [main] Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
2017-12-04 11:41:14,201 WARN [main] mapreduce.ExportJobBase: IOException checking input file header: org.apache.hadoop.ipc.RemoteException(org.apache.hadoop.hdfs.server.namenode.SafeModeException): Zero blockLocations for /sqoopout/part-m-000000. Name node is in safe mode.
The reported blocks 0 needs additional 68 blocks to reach the threshold 0.9999 of total blocks 68.
The number of live datanodes 0 has reached the minimum number 0. Safe mode will be turned off automatically once the thresholds have been reached.
at org.apache.hadoop.hdfs.server.namenode.FSNamesystem.getBlockLocations(FSNamesystem.java:1806)

```

```

mysql> select * from company;
+----+-----+-----+
| id  | name   | location |
+----+-----+-----+
| 1   | acadgild | Bangalore |
| 2   | eduvista | Walnut   |
| 3   | eduvista | Bangalore |
+----+-----+-----+
3 rows in set (0.00 sec)

```